

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import sklearn
6 import seaborn as sns
7 import warnings
8 warnings.filterwarnings('ignore')
9 plt.rcParams["figure.figsize"] = [10,5]
10 # Ignore warnings
11
12 import warnings
13 # Set the warning filter to ignore FutureWarning
14 warnings.simplefilter(action='ignore',category=FutureWarning)
```

## TRAINING DATA PRE-PROCESSING

```
In [6]: 1 df = pd.read_csv('C:\\Users\\Super\\Downloads\\USA_Housing.csv')
        2 df
```

Out[6]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06	208 Michael Ferry Apt. 674\\nLaurabury, NE 3701...
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06	188 Johnson Views Suite 079\\nLake Kathleen, CA...
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06	9127 Elizabeth Stravenue\\nDanielstown, WI 06482...
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06	USS Barnett\\nFPO AP 44820
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05	USNS Raymond\\nFPO AE 09386
...	...	...	...	...	...	...	...
4995	60567.94414	7.830362	6.137356	3.46	22837.36103	1.060194e+06	USNS Williams\\nFPO AP 30153- 7653
4996	78491.27543	6.999135	6.576763	4.02	25616.11549	1.482618e+06	PSC 9258, Box 8489\\nAPO AA 42991-3352
4997	63390.68689	7.250591	4.805081	2.13	33266.14549	1.030730e+06	4215 Tracy Garden Suite 076\\nJoshualand, VA 01...
4998	68001.33124	5.534388	7.130144	5.44	42625.62016	1.198657e+06	USS Wallace\\nFPO AE 73316
4999	65510.58180	5.992305	6.792336	4.07	46501.28380	1.298950e+06	37778 George Ridges Apt. 509\\nEast Holly, NV 2...

5000 rows × 7 columns

## Data Shape

```
In [7]: 1 # Data shape
        2 print('train data:', df.shape)
```

train data: (5000, 7)

In [8]:

```
1 # View first few rows
2 df.head(5)
```

Out[8]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05	USNS Raymond\nFPO AE 09386

In [9]:


```
1 # Data Info
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null  float64
1   Avg. Area House Age                   5000 non-null  float64
2   Avg. Area Number of Rooms             5000 non-null  float64
3   Avg. Area Number of Bedrooms          5000 non-null  float64
4   Area Population                       5000 non-null  float64
5   Price                                5000 non-null  float64
6   Address                              5000 non-null  object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

## Missing Data

```
In [10]: 1 # Heatmap
          2 sns.heatmap(df.isnull(),yticklabels = False, cbar = False,cmap = 'tab20c_r')
          3 plt.title('Missing Data: Training Set')
          4 plt.show()
```

Missing Data: Training Set



```
In [11]: 1 # Remove Address feature
        2 df.drop('Address', axis = 1, inplace = True)
```

```
In [12]: 1 # Remove rows with missing data
        2 df.dropna(inplace = True)
```

```
In [13]: 1 df.head()
```

Out[13]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05

## Numeric Features

```
In [14]: 1 df.describe()
```

Out[14]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562390	5.322283	6.299250	3.140000	29403.928700	9.975771e+05
50%	68804.286405	5.970429	7.002902	4.050000	36199.406690	1.232669e+06
75%	75783.338665	6.650808	7.665871	4.490000	42861.290770	1.471210e+06
max	107701.748400	9.519088	10.759588	6.500000	69621.713380	2.469066e+06

# GETTING MODEL READY

```
In [15]: 1 # Shape of train data  
2 df.shape
```

```
Out[15]: (5000, 6)
```

## OBJECTIVE 2: MACHINE LEARNING

Next, I will feed these features into various classification algorithms to determine the best performance using a simple framework: Split, Fit, Predict, Score It.

### Target Variable Splitting

```
In [16]: 1 #Split data to be used in the models  
2 # Create matrix of features  
3 x = df.drop('Price',axis=1) # grabs everything else but 'Price'  
4  
5 # Create target variable  
6 y = df['Price'] # y is the column we're trying to predict  
7
```

```
In [17]: 1 x.shape
```

```
Out[17]: (5000, 5)
```

```
In [18]: 1 y.shape
```

```
Out[18]: (5000,)
```

```
In [19]: 1 # Use x and y variables to split the training data into train and test set
2 from sklearn.model_selection import train_test_split
3 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = .20, random_state = 101)
```

```
In [20]: 1 x_train.shape
2 x_train
```

Out[20]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
<b>3413</b>	69048.78809	6.619712	6.123813	4.33	36817.36876
<b>1610</b>	67866.89993	5.393978	9.359022	5.44	43122.57418
<b>3459</b>	56636.23819	5.497667	7.121872	6.10	47541.43176
<b>4293</b>	79310.36198	4.247434	7.518204	4.38	43982.18896
<b>1039</b>	72821.24766	6.480819	7.116655	5.33	40594.05930
...	...	...	...	...	...
<b>4171</b>	56610.64256	4.846832	7.558137	3.29	25494.74030
<b>599</b>	70596.85095	6.548274	6.539986	3.10	51614.83014
<b>1361</b>	55621.89910	3.735942	6.868291	2.30	63184.61315
<b>1547</b>	63044.46010	5.935261	5.913454	4.10	32725.27954
<b>4959</b>	75078.79152	7.644779	8.440726	4.33	56148.44932

4000 rows × 5 columns



In [21]:

```
1 y_train.shape
2 y_train
```

Out[21]:

```
3413    1.305210e+06
1610    1.400961e+06
3459    1.048640e+06
4293    1.231157e+06
1039    1.391233e+06
...
4171    7.296417e+05
599     1.599479e+06
1361    1.102641e+06
1547    8.650995e+05
4959    2.108376e+06
Name: Price, Length: 4000, dtype: float64
```

In [22]:

```
1 x_test.shape
2 x_test
```

Out[22]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
<b>1718</b>	66774.99582	5.717143	7.795215	4.32	36788.980330
<b>2511</b>	62184.53937	4.925758	7.427689	6.22	26008.309120
<b>345</b>	73643.05730	6.766853	8.337085	3.34	43152.139580
<b>2521</b>	61909.04144	6.228343	6.593138	4.29	28953.925380
<b>54</b>	72942.70506	4.786222	7.319886	6.41	24377.909050
...	...	...	...	...	...
<b>3900</b>	77615.85134	6.200603	6.909327	2.27	36591.523450
<b>3753</b>	66925.19935	5.153050	8.396903	3.16	42590.685170
<b>3582</b>	71778.02618	5.921280	7.411045	4.00	37634.041320
<b>2392</b>	87272.09339	5.025866	7.184765	5.39	7522.333138
<b>3343</b>	70271.10419	5.856327	6.782116	2.46	28101.644400

1000 rows × 5 columns

# LINEAR REGRESSION

Model Training

```
In [25]: 1 # Fit
          2 # Import model
          3 from sklearn.linear_model import LinearRegression
          4
          5 # Create instance of model
          6 lin_reg = LinearRegression()
          7
          8 # Pass training data into model
          9 lin_reg.fit(x_train, y_train)
```

```
Out[25]: ▾ LinearRegression
          LinearRegression()
```

## Model Testing

Class prediction

```
In [26]: 1 #predict
          2 y_predict=lin_reg.predict(x_test)
          3 y_predict.shape
          4 y_predict
```

```
Out[26]: array([1257919.72924299,  822112.41868756, 1740669.05869474,
                972452.12926804,  993422.2632988 ,  644126.07416935,
                1073911.79097589,  856584.00208537, 1445318.25527738,
                1204342.19071515, 1455792.46233196, 1298556.65691754,
                1735924.33854636, 1336925.77593212, 1387637.43241543,
                1222403.77757898,  613786.28673738,  963933.54403085,
                1221197.33061287, 1198071.57580528,  505861.89541388,
                1769106.54726586, 1853881.16845511, 1200369.50514846,
                1065129.12845899, 1812033.73048156, 1768686.47104264,
                1439920.83823817, 1387251.9966963 , 1541178.39227172,
                726418.80525623, 1754497.609143 , 1462185.72661629,
                1025600.16064332, 1284926.86862687,  917454.59581447,
                1187046.94951786,  999330.91123324, 1329536.63408978,
                782191.60431848, 1393272.03057331,  578216.88372019,
                822643.37151103, 1895533.11423642, 1672019.84904555,
                966926.45430148, 1129674.55621678,  792797.75924288,
                1161057.18404066, 1472396.7143581 , 1457656.70413195,
                1162939.33425471, 1099453.68096241, 1358107.44627459,
                841103.7037299 ,  986322.30559828, 1123323.53002156,
                1252520.62221222,  4122270.66222711,  402102.22221171,
```

```
In [27]: 1 # Combine actual and predicted values side by side
2 results = np.column_stack((y_test, y_predict))
3 print("Actual Values | Predicted values")
4 print("-----")
5 for actual, predicted in results:
6     print(f"{actual:14.2f} | {predicted:12.2f}")
7
```

Actual Values	Predicted values
-----	
1251688.62	1257919.73
873048.32	822112.42
1696977.66	1740669.06
1063964.29	972452.13
948788.28	993422.26
730043.65	644126.07
1166925.15	1073911.79
705444.12	856584.00
1499988.88	1445318.26
1288199.15	1204342.19
1441736.76	1455792.46
1279681.15	1298556.66
1754969.16	1735924.34
1511653.46	1336925.78
1441956.20	1387637.43
1119992.62	1222403.78
727866.53	613786.29
1128885.10	862833.54

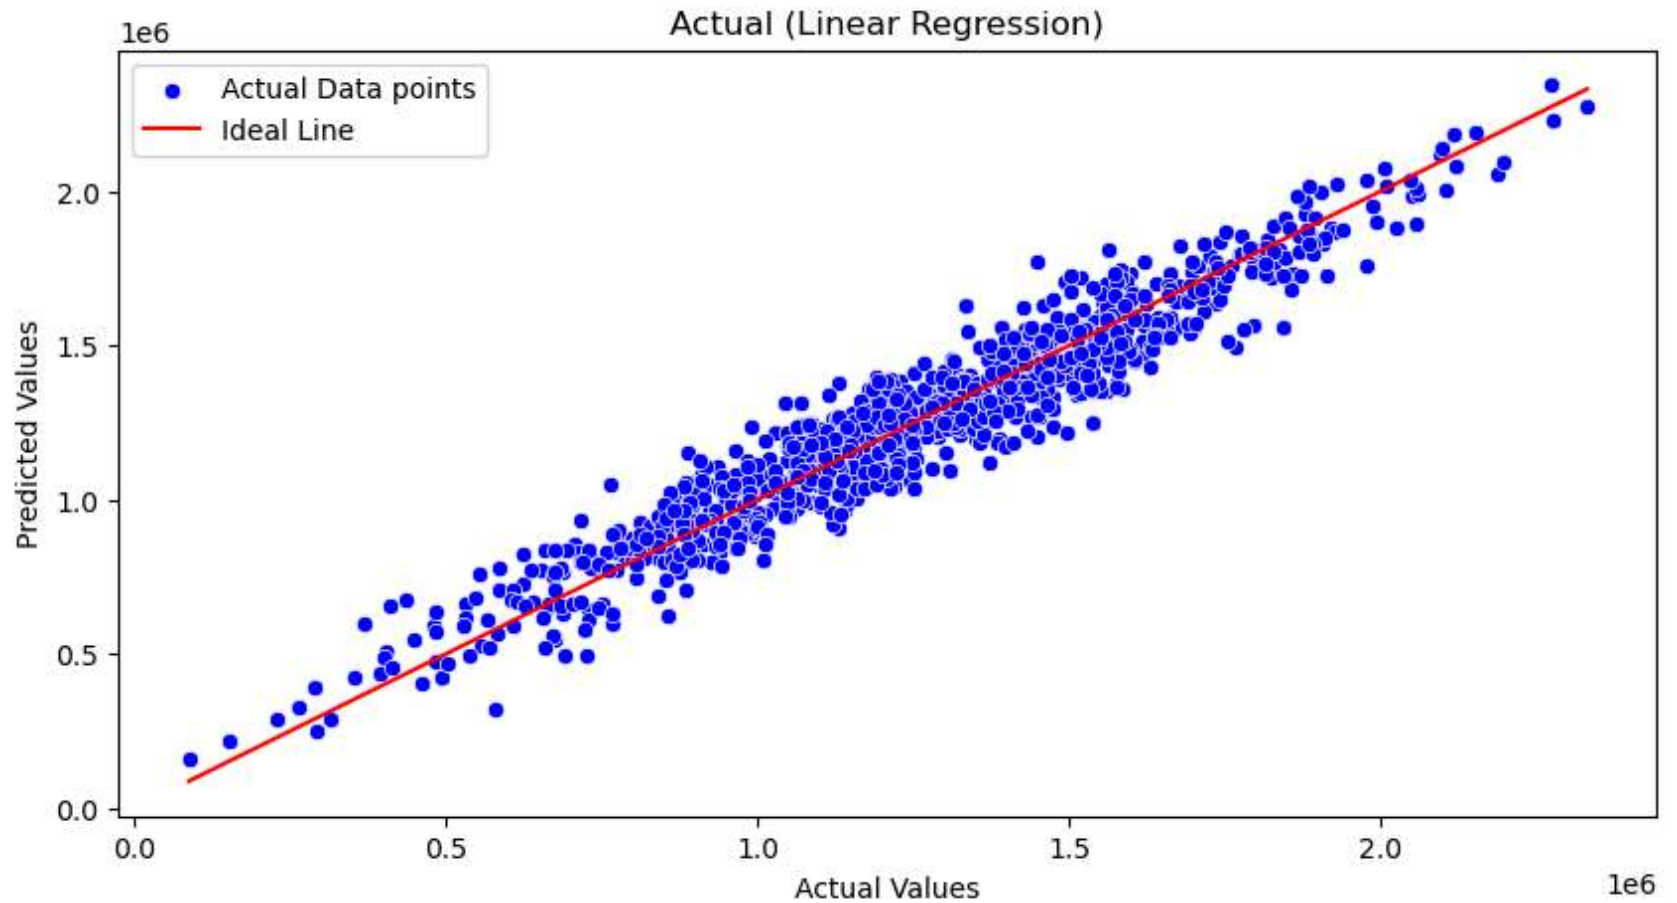
## Residual Analysis

In [28]:

```
1 residual = actual - y_predict.reshape(-1)
2 print(residual)
```


```
[-6.97228472e+04  3.66084463e+05 -5.52472177e+05  2.15744753e+05
 1.94774619e+05  5.44070808e+05  1.14285091e+05  3.31612880e+05
-2.57121373e+05 -1.61453087e+04 -2.67595580e+05 -1.10359775e+05
-5.47727457e+05 -1.48728894e+05 -1.99440550e+05 -3.42068956e+04
 5.74410595e+05  2.24263338e+05 -3.30004486e+04 -9.87469381e+03
 6.82334987e+05 -5.80909665e+05 -6.65684286e+05 -1.21726231e+04
 1.23067754e+05 -6.23836848e+05 -5.80489589e+05 -2.51723956e+05
-1.99055115e+05 -3.52981510e+05  4.61778077e+05 -5.66300727e+05
-2.73988845e+05  1.62596721e+05 -9.67299866e+04  2.70742286e+05
 1.14993248e+03  1.88865971e+05 -1.41339752e+05  4.06005278e+05
-2.05075149e+05  6.09979998e+05  3.65553510e+05 -7.07336232e+05
-4.83822967e+05  2.21270428e+05  5.85223258e+04  3.95399123e+05
 2.71396980e+04 -2.84199832e+05 -2.69459822e+05  2.52575477e+04
 8.87432010e+04 -1.69910564e+05  3.47093178e+05  2.01874576e+05
 6.48733520e+04 -6.53417503e+04 -2.40082781e+05  6.89093673e+05
-2.74620201e+05  7.94525134e+04  5.27608365e+05 -5.88347159e+04
-1.54547673e+05 -1.69309109e+05  3.69338470e+05 -2.98971867e+05
-2.08050774e+05  3.03118502e+05  3.30651051e+05 -2.46195692e+04
 8.71276104e+04 -6.46394947e+05  2.64173916e+05  3.32850074e+05
 2.66126122e+05  2.00623332e+05  4.10063342e+05  4.10070102e+05]
```

```
In [29]: 1 sns.scatterplot(x=y_test, y=y_predict, color='blue', label='Actual Data points')
2 plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', label='Ideal Line')
3 plt.xlabel('Actual Values')
4 plt.ylabel('Predicted Values')
5 plt.title('Actual (Linear Regression)')
6 plt.legend()
7 plt.show()
```



```
In [30]: 1 # Distribution plot for Residual (difference between actual and predicted values)
         2 sns.distplot(residual, kde=True)
```

```
Out[30]: <Axes: ylabel='Density'>
```



## Model Evaluation

```
In [31]: 1 # Score It
2 from sklearn.metrics import mean_squared_error
3
4 print('Linear Regression Model')
5 # Results
6 print('--'*30)
7 # mean_squared_error(y_test, y_pred)
8 mse = mean_squared_error(y_test, y_predict)
9 rmse = np.sqrt(mse)
10
11 # Print evaluation metrics
12 print("Mean Squared Error:", mse)
13 print("Root Mean Squared Error:", rmse)
```

Linear Regression Model

-----

Mean Squared Error: 10100187856.996004

Root Mean Squared Error: 100499.69083035034

```
In [ ]:
```

```
1
```